

CS 477/677 Causal Inference: Homework 1

Probability and Statistics

Due: Tuesday September 20, 2016, 11:59pm

40 (50) Points Total

Ilya Shpitser
ilyas@cs.jhu.edu

1 Introduction

Homeworks in this class will consist of a programming part and a conceptual part.

1.1 Programming Assignments and the Background on the R Language

The programming part of the homework is meant to familiarize you with tools needed to perform causal inferences from data, and give you a chance to practice doing such analyses. This part will involve writing and running code, and perhaps answering some questions about the results. You will implement all assignments using the R programming language. Submissions in other languages will not be accepted. R is a very powerful high level programming language for data analysis and statistical computing. While no programming language is perfect, R's expressiveness and built in functions for common statistical operations will make writing data analysis programs simple compared to more "heavy duty" general purpose languages such as Java. R's semantics resemble the LISP programming language in many ways, while R's syntax resembles standard syntax used by the Algol family of languages (such as C, C++, Java, and python). If you are familiar with one of these languages, R will not be difficult for you to pick up and use. R can be downloaded for free at <https://www.r-project.org/>.

The R language has many predefined functions and a very large number of contributed packages for all sorts of things. To help you navigate, we will generally give suggestions for which functions or packages might be relevant. R is an interpreted programming language. You can evaluate code directly using a "read eval print loop" interpreter, load R scripts into an interpreter and run it there, or run R scripts from the command line using the command `Rscript` that comes with R. The most important R function for a beginner is `help(.)`. You can use it get a detailed description of any built in function in R, and generally on functions that are implemented in packages you download (provided those packages are well-documented!)

1.2 Conceptual Assignments and the Background on Latex

The conceptual part of the homework is meant to help you think through the philosophical and mathematical issues involved in causal inference. This part may involve manual calculation, mathematical and logical arguments, perhaps even essay questions. Your answers to the conceptual part of the homework must be submitted as a pdf file generated by latex. Latex is a very powerful typesetting system particularly useful for writing scientific and mathematical documents. While latex involves a bit of a learning curve

compared to “what you see is what you get” editors such as Microsoft Word, the time investment is generally worth it. If you want to write scientific publications, latex is always a safe choice, as many journals and peer reviewed conferences have standardized on it. Latex can be viewed as a programming language that is compiled into a pdf file. A latex source file contains a combination of plain text and directives that instruct the compiler how to typeset the document. We recommend you use the command `pdflatex` to convert latex source code into pdf. This command is generally available on most common operating systems in use today. We will provide the latex source code for every homework assignment that we recommend you use as a standard point for the conceptual part of the assignment.

2 Programming (20 Points)

In this part of the assignment, you will familiarize yourself with the R language, and write functions for creating data frames (R data structures used to represent data sets), performing non-parametric bootstrap, constructing confidence intervals for any statistic of interest, and fitting parameters in a linear regression model using maximum likelihood. Unless otherwise stated, please implement all code yourself (even if an existing implementation already exists as a package, or even a part of standard R!)

2.1 Linear Regression (6 Points)

You will implement maximum likelihood fitting of a linear regression using a data set included with the homework. The following description of the dataset appears in [1], Chapter 8:

This dataset is from the Job Search Intervention Study (JOBS II) [2]. In the JOBS II field experiment, 1,801 unemployed workers received a pre-screening questionnaire and were then randomly assigned to treatment and control groups. Those in the treatment group participated in job-skills workshops. Those in the control condition received a booklet describing job-search tips. In follow-up interviews, two key outcome variables were measured: a continuous measure of depressive symptoms based on the Hopkins Symptom Checklist (`depress2`), and a binary variable representing whether the respondent had become employed (`work1`). In the JOBS II data, a continuous measure of job-search self-efficacy represents a key mediating variable (`job_seek`). In addition to the outcome and mediators, the JOBS II data also include the following list of baseline covariates that were measured prior to the administration of the treatment: pretreatment level of depression (`depress1`), education (`educ`), income, race (`nonwhite`), marital status (`marital`), age, sex, previous occupation (`occp`), and the level of economic hardship (`econ_hard`).

In this assignment, we will use a subset of the original dataset for which complete data on all features and outcome is available. We will discuss how to handle missing data later in the course. Your regression model will regress an intercept (represented by a column of 1s) and variables `treat`, `econ_hard`, `depress1`, `sex`, `age`, `occp`, `marital`, `nonwhite`, `educ`, `income`, and `job_seek` (the first 11 columns of the dataset) with the outcome variable `depress2` (the 12th column of the dataset). The R language has functions which implement maximum likelihood fitting for linear regression models (`lm` and

glm). Please do not use them in this assignment, and implement fitting directly using matrix manipulation. You may find the matrix inversion function `solve(.)` and the matrix multiplication operator `%*%` helpful.

```
# Input:
# Output: A row vector of linear regression coefficients, for a model with an
#          intercept and input features corresponding to the first 11 columns, and
#          the output the 12th column of a data matrix obtained from
#          a file "Jobs-NoMiss-Cont.tab".
linreg = function(){

  d <- read.table("Jobs-NoMiss-Cont.tab", header = TRUE)
  m <- as.matrix(d)

}
```

2.2 The Bootstrap

2.2.1 Implementation (10 Points)

You will implement the following four functions.

```
# Input:
#          mat: a matrix with n rows and k columns (rows are samples, columns are
#          variables).
# Output: a data frame with column (variable) names "x1" to "xk", and data from the
#          matrix.
df_make = function(mat){

}
```

This function merely wraps a data frame around a matrix representing a data set, and gives standard names to columns (representing features or variables in our data). You might find functions `matrix(.)`, `as.data.frame(.)`, `dim(.)`, `paste(.)` and `colnames(.)` helpful.

```
# Input:
#          df: a data frame with n rows to be resampled.
# Output: a data frame with n rows obtained from the input dataframe
#          by sampling rows with replacement.
df_resample = function(df){

}
```

This function constructs an appropriately resampled dataset used in nonparametric bootstrap. You might find the function `sample(.)`, and reading up on R's `1:k` syntactic sugar notation for generating vectors helpful.

```
# Input:
#          df: a data frame to be resampled
#          k: number of resampled datasets to generate.
#          f: a function of df giving the bootstrap statistic
#             (e.g. function(df) { mean(df$x1) })
#          q: a real number 0 < q < 0.5 giving the lower quantile of the desired
#             confidence interval.
```

```
# Output:  a three element vector giving the statistic of interest (first element),
#           and lower and upper confidence intervals corresponding to
#           q and 1-q quantiles (second and third elements) of the empirical
#           bootstrap distribution.
bootstrap_ci = function(df, k, f, q){
}

```

This function uses non-parametric bootstrap to construct confidence intervals around a parameter of interest. We will use a particular rule which estimates this parameter using a data set. For example, if we were interested in the mean of the variable X_1 , and we had data frame with a set of values for X_1 , then we could use the maximum likelihood procedure for the mean. We will implement these rules by a *reified function*, sometimes known as a *closure* passed as the third argument. A closure that calculates the mean of X_1 from a data frame `df` could be constructed as follows:

```
f = function(df) { mean(df$x1) }
```

This creates a variable `f` in R which is a function of one argument, that expects a data frame, and which calculates and returns the mean of the `x1` column of this data frame. This variable can be treated as an ordinary variable and passed around to other functions, and it can be evaluated at any point by calling the variable with an argument: `f(df)`.

The function you will implement will first evaluate `f` on the original data set to obtain the statistic, generate `k` resampled data sets using the function `df_resample` above, store the values of the statistic of interest for each of these resampled datasets, and obtain the appropriate confidence intervals using the `quantile(.)` function applied to the difference of the original statistic, and the vector of statistics obtained from the resampled datasets, corresponding to the empirical bootstrap distribution. Remember, reported confidence intervals are centered around the statistic of interest. For example, if the statistic is an estimated mean of 3.2, the 0.05 quantile is -0.72 , and the 0.95 quantile is 0.68 , we report the vector 3.2, 2.48, 3.88.

```
# Input:
#       none
# Output:
#       none
main = function(){
    set.seed(0)
}

```

This function will always be the entry point into your homework code. R does not require that the entry point be named `main` specifically, but we still use this name as a nod to other languages where `main` was a special function. We will always start this function with a call to `set.seed(0)`. This call sets the seed of the pseudo-random number generators used by R. We will always use this seed in order to make everyone's code behave in the same way, and eliminate randomness from the grading process.

In this function, you will generate a 1000 sample dataset using a multivariate Gaussian

distribution with parameters

$$\mu = (1, 2, 3, 0, 0), \Sigma = \begin{pmatrix} 1, 1, 1, 1, 1 \\ 1, 3, 1, 1, 1 \\ 1, 1, 5, 1, 1 \\ 1, 1, 1, 2, 0 \\ 1, 1, 1, 0, 2 \end{pmatrix},$$

and obtain the 95% confidence intervals (using the 0.025 and 0.975 quantiles) using this data set and the function `bootstrap_ci(.)` for the following statistics: the mean of X_1 , the mean of X_2 , the median of X_3 , and the covariance of X_4 and X_5 . This function will return nothing, but will print the output of `confidence_ci(.)` (in the order listed here), followed by the output of `linreg()` using the `print(.)` function. You might find functions `mean(.)`, `median(.)`, `var(.)` and `cov(.)` helpful. Use the function `mvnrm(.)` to sample from a multivariate Gaussian. This function is found in a package MASS, which you can install using the `install.packages(.)` function in the R interpreter, and import into your code using the command `library(MASS)`.

```
main()
```

Place this call at the end of your file such that running it as a script will evaluate the entry point function.

2.2.2 Questions (4 Points)

- 1 Which confidence interval is wider, for the mean of X_1 or the mean of X_2 . Why do you think that is?
- 2 What substantive conclusions can you draw from the coefficient values of the regression model you fit? Depression was evaluated by a questionnaire where a lower value is better (less depressed).

3 Analytical (20 Points)

1 (3 points) Reichenbach's Common Cause Principle states that if two variables A and B are associated, then one of three things is true, either

- 1 A causes B (directly or indirectly),
- 2 B causes A (directly or indirectly),
- 3 there is a common cause of both A and B .

Give an example of a pair of associated A and B that satisfies condition 1, another pair of associated A and B that satisfies condition 3, and another pair of associated A and B that satisfies none of the above three conditions.

2 (4 points) Consider two versions of Pearl's example. We record the recovery rates of 700 patients who were given access to the drug. A total of 350 patients chose to take the drug and 350 patients did not. The results are shown below.

Version 1: men and women have different rates:

	Drug	Placebo
Men	81 out of 87 recovered (93%)	234 out of 270 recovered (87%)
Women	192 out of 263 recovered (73%)	55 out of 80 recovered (69%)
Total	273 out of 360 recovered (78%)	289 out of 350 recovered (83%)

Version 2: some patients complained of a side effect (even under placebo), and recovery rates seem different in those who complained:

	Drug	Placebo
No side effect	81 out of 87 recovered (93%)	234 out of 270 recovered (87%)
Side effect	192 out of 263 recovered (73%)	55 out of 80 recovered (69%)
Total	273 out of 360 recovered (78%)	289 out of 350 recovered (83%)

Would you give the drug in version 1 of the problem? In version 2 of the problem? Explain.

3 (2 points) Why do you think Simpson's paradox is a paradox?

4 (6 points) An *exponential distribution* for a continuous non-negative random variable has the density function of the form

$$f(X) = \lambda \exp(-\lambda X).$$

- Assume we have n independent samples x_1, \dots, x_n from an exponential distribution with an unknown parameter λ . Write down the log likelihood, and derive the maximum likelihood estimate $\hat{\lambda}$.
- Derive $E[X]$ if X is a random variable with an exponential distribution with parameter λ . Hint: use the empirical approximation of the mean.

5 (5 points) Show that conditional independence satisfies the semi-graphoid axioms.

4 Extra Credit (10 points)

1 (7 points) Show that if $p(C, A, Y)$ is a distribution, then for any value a of A ,

$$q_a(Y, C) = p(Y \mid A = a, C)p(C)$$

is also a distribution.

2 (3 points) Simpson's paradox occurs in joint distributions $p(Y, A, C)$ where

$$p(Y \mid A = a_1) < p(Y \mid A = a_2), \text{ but } p(Y \mid A = a_1, C = c) > p(Y \mid A = a_2, C = c)$$

for any value of c of C . Define $q_a(Y, C)$ as in question 1, and let $q_a(Y) = \sum_C q_a(Y, C)$, $q_a(Y \mid C) = q_a(Y, C)/q_a(C)$. Show that if for any value c of C , $q_{a_1}(Y \mid C = c) < q_{a_2}(Y \mid C = c)$, then $q_{a_1}(Y) < q_{a_2}(Y)$.

5 How to Submit

You can obtain the assignment from either the class piazza page or on the class blackboard page. You will submit the assignment on blackboard. There will generally be no extensions to the listed deadline.

6 What to Submit

A pdf file named homework1.pdf, and an R script named homework1.R. Please include your name, and email at the top of both files.

7 Questions?

You can ask questions about the homework, or any other aspect of the class at the class piazza page, found here:

<http://piazza.com/jhu/fall2016/cs600477677/home>

To signup for piazza for the class, follow this link:

<http://piazza.com/jhu/fall2016/cs600477677>

References

- [1] Hrishikesh D. Vinod. *Advances in Social Science Research Using R*. Springer-Verlag New York, 2010.
- [2] A. Vinokur and Y. Schul. Mastery and inoculation against setbacks as active ingredients in the jobs intervention for the unemployed. *Journal of Consulting and Clinical Psychology*, 65(5):867–877, 1997.